



EXERCISES

Course	Digital Image Processing & Computer Vision
Term	February 2021
Main Textbook	Digital Image Processing (4th Edition), Rafael C. Gonzalez
Goal	Practically solve computer vision and image processing problems
Score	6-7 (Vary for bachelor and master students)
Number of HW	6 (+1 as an extra homework to improve the score)

Instructions The reports should be written in the considered template and include the four sections as:

- 1- Abstract: 1 paragraph
- 2- Technical description: 1 to 2 pages
Describe the idea and solution for solving problems in detail.
- 3- Describe results: 3 to 4 pages
Display and discuss the experimental results including table, images, etc.
- 4- Appendix: Code with **comments**, Reference, etc.

Notice: Depending on the type of exercise, the number of pages may be increased.

Delivery	<ul style="list-style-type: none">• All homework should be sent through VU (No Telegram, Email, etc.).• The report in PDF format named "Number of Homework-First Name Last Name.pdf".• Notice the deadlines.
Points	<ul style="list-style-type: none">• Don't use MATLAB, Python, etc. library/toolbox for solving problems (Except math functions).• Utilize functions when explicitly mentioned in the question.• You can compare your own result by MATLAB or etc. output (optional).• Any form of plagiarism will not be entertained and will result in a loss of grade.• The symbol * at the beginning of some question represents the optionality implementation.• Never take screenshots of generated images.• Try to clearly answer the question e.g., 1.1.1., 1.1.2, etc.• <u>Discuss and comment on the obtained results.</u>

BE HAPPY 😊

Behrouz Bolourian Haghighi
Lecturer in Computer Engineering
Senior Researcher at Machine Vision Laboratory
Ph.D. Candidate in Artificial Intelligence and Robotics
Computer Engineering Department, Ferdowsi University of Mashhad, Iran
Homepage: <http://b.bolourian.student.um.ac.ir/>
Laboratory: <http://mvlab.um.ac.ir/>

1. Image Fundamentals

1.1. Geometric Transformation

- 1.1.1. *You are given two pictures of the same scene, taken at different times. To align the two pictures, you need to find a mapping function between the two pictures based on some common feature points. Suppose you were able to extract N ($N \geq 3$) feature points in both images that correspond to the same set of object features, with image coordinates given as (u_k, v_k) and (x_k, y_k) , $k = 1, 2, \dots, N$. Also, suppose you want to use an affine mapping to approximate the actual unknown mapping function. How would you determine the affine mapping parameters?
- 1.1.2. Suppose you want to make a panoramic picture of a wide landscape out of two separately captured pictures with a certain overlap. Propose and implement an algorithm for stitching up the two images to make a panorama. List the steps involved; Also, discuss and display the results (Test image Car1&2).
- 1.1.3. Write a program that implements rotation (with interpolation) of an image by angles (45° , 100° , 670°), and apply it to Barbara image. The rotation center should be the image center.

1.2. Quantization & Interpolation

- 1.2.1. For two cases as without and with histogram equalization (uniform histogram), display the quantized image in (4, 8, 16, 32, 64, 128) Levels and its histograms. Also, the optimum mean square error obtained for each case. Discuss and report the results for the gray Elaine image. It should be noted, you can use `rgb2gray`, `histeq` and `immse` functions for this problem.

Report MSE						
Level	4	8	16	32	64	128
Without Histeq						
With Histeq						

- 1.2.2. Write a program which can, firstly, downsample an image by a factor of 2, with and without using the averaging filter, and also, up-sample the previously downsampled images by a factor of 2, using the pixel replication and bilinear interpolation methods, respectively. Display (zoom of image) and discuss the results obtained with different methods for the Goldhill image. Note, you can use `immse` function for this problem.

Report MSE		
	Pixel Replication	Bilinear Interpolation
Averaging		
Remove Row&Column		

- 1.2.3. The initial image consists of eight bits of data for each pixel. Create new images using 5, 4, 3, 2 and 1 bit only for each pixel. How many bits are needed to preserve image quality? Does it change from place to place in the image? Discuss about the results. (Test Image Grayscale of Barbara).

2. Contrast Adjustment

2.1. Histogram Equalization

- 2.1.1. Write a program that can compute the histogram of a grayscale image (assuming 256 levels of gray). In a separate main program, apply the program to Camera Man image, and illustrate the histogram as a stem plot besides the test image (using "subplot" function).
 - 2.1.1.1. Decrease the brightness of Camera Man by dividing the intensity values by 3 and named output as D.
 - 2.1.1.2. Plot the histograms of Input and D. What can you observe from these two histograms?
 - 2.1.1.3. Perform histogram equalization on D and output the result as H.
 - 2.1.1.4. Perform local histogram equalization on image D and output the result as L.
 - 2.1.1.5. Plot the histograms of H and L. What's the main difference between local and global histogram equalization?

2.1.2. Write a program that performs histogram equalization on Camera Man image. Display the original and equalized images, as well as their corresponding histograms, all in one figure as mentioned in 2.1.1.

2.1.3. What is the difference between histeq and imadjust functions in Matlab? Play with these functions with various input parameters for Camera Man image. Write down your observations in your report and display results.

2.2.1. Implement a local histogram equalization with different windows size for the HE1,2,3, and 4 images. Explain and display the results. Discuss the effects of increasing window size and compare it with global histogram equalization in detail.

- 3.1.1. *Why are box filters bad smoothing filters? List all reasons.
- 3.1.2. *Do the bad features improve if apply the filters several times?
- 3.1.3. What is the resulting filter if apply the 3×3 box filter many times? (Test on grayscale Elaine Image).
- 3.1.4. How does the size of the mask affect blurring and noise reduction? (Test on grayscale Elaine Image).
- 3.1.5. Which mask size do you think provides a better tradeoff between blurring and noise reduction for this image?
- 3.1.6. What is the resulting if apply Laplacian mask ($\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$) many times? (Test output of 3.1.3).

3.2.1. Write a program that can, first, add salt-and-pepper noise to an image with a specified noise density. Try different noise density (0.05, 0.1, 0.2, 0.4). Then, perform median filtering with a specified window size. Consider only the median filter with a square shape. For each density, discuss the effect of filtering with different window sizes (3, 5, 7, 9) and experimentally determine the best window size. Note: you can use imnoise and immse functions to generate noisy images and compare the quality of images, respectively. Also, you can ignore the boundary problem by only performing the filtering for the pixels inside the boundary. (Test on grayscale Elaine Image).

	Report MSE				
	3×3	5×5	7×7	9×9	11×11
$\rho = 0.05$	0.0000	0.0000	0.0000	0.0000	0.0000
$\rho = 0.1$	0.0000	0.0000	0.0000	0.0000	0.0000
$\rho = 0.2$	0.0000	0.0000	0.0000	0.0000	0.0000
$\rho = 0.5$	0.0000	0.0000	0.0000	0.0000	0.0000

	Report MSE									
	Median					Box Filter				
	3×3	5×5	7×7	9×9	11×11	3×3	5×5	7×7	9×9	11×11
$\sigma = 0.01$										
$\sigma = 0.05$										
$\sigma = 0.1$										

3.2.3. Design proper filters to simultaneously remove Gaussian and salt & pepper noise. Please detail the steps of the denoising process and specify corresponding parameters. Provide some discussions about the reason why those filters and parameters are chosen.

3.3. Sharpening, Blurring, and Noise Removal

3.3.1. Take blurry and noisy images (shooting in low light is a good way to get both) by your cellphone and try to improve their appearance and legibility. Display and discuss the results before and after improving.

3.4. Edge Detection

3.4.1. These are often used first-order difference filters in x-direction:

$$a) \frac{1}{2} \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}, \quad b) \frac{1}{6} \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}, \quad c) \frac{1}{8} \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Compare and describe the properties of the three filters; In the following discuss the efficiency of filters for edge detection problems. Are these filters suitable for computing the 2-D gradient? (Test on Elaine Lena Image).

3.4.2. Robert suggested the filter to compute the 2-D gradient and to detect edges.

$$a) \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}, \quad b) \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

First, in which directions do these filters detect edges? Second, compare the quality of this filter with the filter from previous steps. (Test on grayscale Elaine Image).

3.5. Unsharp Masking

3.5.1. A simple unsharp masking filter has the following form:

$$(1 - \alpha)I + \alpha I' = I + \alpha(I' - I),$$

where $\alpha \in [0,1]$ represents the threshold step of strength sharpening and I' is smoothed version of I which obtained using different filter sizes based on the gaussian filter. Compare and discuss the results for different filter sizes (3,5,7,9,11) and α for the grayscale Lena image. What happens if we set α very large or small? How to obtain the optimum value? Note: You can use imgaussfilt function for smoothing.

4. Frequency Domain

4.1. Fourier transform

4.1.1. For each filter given below, compute its Fourier transform, and illustrate its magnitude response. Determine what is its function (smoothing, edge enhancement or edge detection) based on the filter coefficients as well as its frequency response. For each filter, determine whether it is separable? If yes, compute the FT separately and explain the function of each 1D filter. If not, compute the FT directly. (Test on grayscale Lena Image).

$$a) \frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \quad b) \begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix} \quad c) \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

4.1.2. Perform 2D DFT on grayscale Lena, Barbara, F16, and Baboon images. Display the magnitude of the DFT image with and without shifting and with and without logarithmic. Display and discuss the results. Also, examine in which frequency range the DFT coefficients have large magnitudes and explain why?

4.2. Filtering

4.2.1. *Use DFT function to compute the linear convolution of an image $F(m,n)$ with a filter $H(m,n)$. Let the convolved image be denoted by $Y(m,n)$. Firstly, suppose the image size is 256×256 and the filter size is 11×11 ; What is the required size of the DFT to obtain the convolution of these two? Explain the exact steps to obtain the convolution result. Secondly, suppose we use a 256×256 point DFT algorithm for $F(m,n)$ and $H(m,n)$, and obtain $Z(m,n)$ as $Z = \text{IDFT}(\text{DFT}(X) * \text{DFT}(H))$. The DFT and IDFT in this equation are both 256×256 points. For what values of (m,n) does $Z(m,n)$ equal $Y(m,n)$?

4.2.2. Write a program that filters grayscale Barbara image by zeroing out certain DFT coefficients.

The program consists of three steps:

1. Performing 2D DFT.
2. Zeroing out the coefficients at certain frequencies (see below).
3. Performing inverse DFT to get back a filtered image.

Note: Truncate or scale the image properly such that its range is between 0 and 255.

For part 2, try the following two types of filters:

- a. Let $F(k,l) = 0$ for $TN < \{k,l\} < (1 - T)N$, $T = 1/4, 1/8$ (low-pass filtering).
- b. Let $F(k,l) = 0$ for the following regions:
 - i. $0 \leq \{k \text{ and } l\} \leq TN$;
 - ii. $0 \leq k \leq TN, \text{ and } (1 - T)N \leq l \leq N - 1$;
 - iii. $(1 - T)N \leq k \leq N - 1 \text{ and } 0 \leq \{l\} \leq TN$;
 - iv. $(1 - T)N \leq k \text{ and } l \leq N - 1; T = 1/4, 1/8$

Display and compare the original and processed images. Discuss the function of the two types of filters.

Note: you can use fft2, ifft2, fftshift, and rgb2gray functions for problem 4.

5. Color

5.1. Color space

- 5.1.1. Convert Lena to HSI format, and display the HIS components as separate grayscale images. Observe these images to comment on what does each of the H, S, I components represent. The HSI images should be saved in double precision.
- 5.1.2. *Present and discuss new color space (at least three) in detail which was not introduced in class (Application, Equation, etc.).

5.2. Quantization

5.2.1. Implement uniform quantization of a color image. Your program should do the following:

1. Read a grayscale image into an array.
2. Quantize and save the quantized image in a different array.
3. Compute the MSE and PSNR between the original and quantized images.
4. Display and print the quantized image.

Notice, your program should assume the input values are in the range of (0,256), but allow you to vary the reconstruction level. Record the MSE and PSNR obtained with $L = 64, 32, 16, 8$ and display the quantized images with corresponding L values. Comment on the image quality as you vary L . (Test on Lena Image).

- 5.2.2. For the Lena image, quantize the R, G, and B components to 3, 3, and 2 bits, respectively, using a uniform quantizer. Display the original and quantized color image. Comment on the difference in color accuracy.
- 5.2.3. We want to weave the Baboon image on a rug. To do so, we need to reduce the number of colors in the image with minimal visual quality loss. If we can have 32, 16 and 8 different colors in the weaving process, reduce the color of the image to these three special modes. Discuss and display the results.

Note: you can use immse and psnr for problem 5.2.

6. Wavelet

6.1. Pyramid

- 6.1.1. For the "Mona Lisa" image, build a 5 level Gaussian pyramid and display it in a format. Also, implement and display a Laplacian (difference of Gaussian (DoG)) pyramid.

- 6.1.2. Describe how separability and cascading can help to speed up Gaussian smoothing and design a fast algorithm for computing a 3-step gaussian pyramid (filtered with $\sigma, \sqrt{2}\sigma, 2\sigma$) of a 2D image using pseudo-code.
- 6.1.3. Given an image of size $N \times N$, where $N = 2^J$, what is the maximum number of levels you can have in an approximation pyramid representation? (The maximum level is reached when the coarsest level has only 1 pixel). What is the total number of pixels in the pyramid (i.e. including pixels at all pyramid levels)? How does this number compare with the original number of pixels in the image? Since this number is larger than the original pixel number, what are some of the benefits of using the approximation pyramid? (give some examples). Repeat the step for the prediction residual pyramid. Display and discuss the results.
- 6.1.4. For the grayscale Lena image, manually compute a 3-level approximation pyramid and corresponding prediction residual pyramid. Use 2x2 averaging for the approximation and use pixel replication for the interpolation filters.
- 6.1.5. For the grayscale Lena Image, compute the wavelet transform (with 3-level) using the Haar analysis filters. Comment on the differences between the pyramids generated in Prob. 6.1.2 with the ones generated here.
- 6.1.6. Quantize all the wavelet coefficients (whole sub-bands) created in Prob. 6.1.3 by a step size of $\gamma = 2$. Then reconstruct the image from the quantized wavelet coefficients using Haar synthesis filter. Report PSNR values and discuss the results.

$$c'(u, v) = \gamma \times \text{sgn}[c(u, v)] \times \text{floor}\left[\frac{|c(u, v)|}{\gamma}\right], c \text{ represents the wavelet coefficient}$$

Note: you can use dwt2, idwt2, and psnr functions for problems 6.

6.2. Denoising

Implement at least two of the various image denoising techniques presented based on wavelet and compare them on both synthetically noised image sequences and real-world (low-light) sequences. Does the performance of the algorithm depend on the correct choice of noise level estimate? Can you draw any conclusions as to which techniques work better? Note: You can freely use any function, toolbox or library. (Test on Lena Image).

7. Features

7.1. Estimate Geometry

- 7.1.1. The geometry of the attacked₂ image (without image filtering-Folder Attack #2) should be reconstructed based on the halftone version as a reference of the image and the attacked₁ image (with image filters-Folder Attack #1). For this aim, there is not any limitation for using available algorithms, functions, and toolbox. After estimating the geometry of the candidate image, the transform matrix should be applied to the attacked₂ image. Finally, the measures including SSIM and MSE between original and reconstructed image should be reported. Also, the number of matched points (MP) between images should be reported for each case.

Type 1 (Histeq)			Type 2 (Sharpen)			Type 3 (Gaussfilt)			Type 4 (Bilatfilt)			Mean			STD		
SSIM	MSE	MP	SSIM	MSE	MP	SSIM	MSE	MP	SSIM	MSE	MP	SSIM	MSE	MP	SSIM	MSE	MP
Result																	

7.2. Corner Detection

- 7.2.1. Implement the Harris corner detector algorithm on the image Building. Superimpose the corner points on to the relevant pixels in the image. It is suggested that you choose a suitable threshold to pick up the interest points, followed by a non-maximal suppression step.